



PRIVATE AI FOR FINTECH

Build vs managed: The private AI reference stack.

Swipe →





THE CORE QUESTION



Can your data legally touch a public model?

- Customer financial data often cannot leave your VPC
- Managed API providers have unpredictable data retention
- Fine-tuning on sensitive data creates regulatory exposure
- The decision is legal and contractual before it is technical





PRIVATE INFERENCE CASE



When private is non-negotiable.

- PII-adjacent financial data in the inference context
- Regulated customer data subject to GLBA or CCPA
- Contractual data residency requirements with customers
- Audit trail must show data never left your infrastructure





MANAGED API CASE



When managed is the right call.

- Non-sensitive internal workflows: drafting, summarization
- Data that is already public or aggregated
- Speed to value outweighs total cost of ownership
- No contractual data residency requirement in scope





THE REFERENCE STACK



Private inference architecture.

- Self-hosted model on GPU cluster in private VPC
- Inference API with no external egress on model requests
- Audit logging at the inference layer, not just application
- Key management and encryption by your infra team





COST REALITY



Private inference **costs more.**

- GPU cluster cost: \$15K-80K/month depending on usage
- Infra engineering: 1-2 dedicated engineers minimum
- Model updates require internal fine-tuning pipeline
- Payback requires significant inference volume to justify





HYBRID ARCHITECTURE



Most fintechs need both.

- Private inference for regulated customer-data workloads
- Managed API for internal tools on non-sensitive data
- Data classification layer to route workloads correctly
- Single orchestration layer spans both inference environments





WHEN TO BUILD



The signals that justify private.

- Inference volume above 10M tokens/month: cost justifies
- Data residency requirement in any customer contract
- Regulatory exam has asked about AI data handling
- Fine-tuning on proprietary data required for accuracy





CODISTE AI

Design your private AI inference stack.

Talk to us ↗

codiste.com